

OntoGuard Layer Facts Brief

Public-safe explanation of the evolved proof surface

Protected by U.S. Patent Application 19/444,521 - Track I prioritized examination granted May 2026.

Layer 1 - AI Output Risk Review

OntoGuard evaluates high-stakes AI outputs, flags support/citation risk, uncertainty, hallucination posture, risk, benchmark posture, evidence gaps, and human-review need without requiring customer documents for first value.

Layer 2 - Decision Authorization Proof

OntoGuard determines whether a proposed AI-driven movement should be ALLOW, BLOCK, or ESCALATE before release. In the sample packet, autonomous release is withheld and routed to human review.

Layer 3 - Execution Boundary and No-Bind Proof

The packet records attempted movement, gate invocation, failed conditions, refused operation, authorization-token status, downstream commit status, protected-effect status, and no-bind receipt.

Layer 4 - Route Completeness Evidence

The current public sample reaches controlled-harness route enforcement: route registered, tested, replayable, fail-closed, controlled non-bypassability asserted, with production enforcement explicitly not asserted.

Layer 5 - Semantic Projection and Learning Signals

Decision projection, coherent governance loading, evidence geodesic, observer views, and review-gated improvement signals explain why the governance state was and how the event can improve future AI behavior.

Boundary statement

The proof surface is public. The protected runtime remains protected. Controlled harness proof is shown; production non-bypassability is not asserted without production integration evidence.