

OntoGuard Layer Facts Brief v4

Protected by U.S. Patent Application 19/444,521 - Track I prioritized examination granted May 2026.

Public-safe explanation of the current proof surface

Layer 1 - AI Output Risk Scan

OntoGuard evaluates representative prompts, responses, logs, and intended actions to identify risk, uncertainty, unsupported claims, review burden, autonomy readiness, and proof gaps without requiring customer documents for first value.

Layer 2 - Runtime Decision Authorization

OntoGuard determines whether a proposed AI-driven state transition should ALLOW, BLOCK, or ESCALATE before consequence. The representative single event escalated to human review and withheld release.

Layer 3 - Execution Boundary and No-Bind Proof

The packet records attempted movement, control point, failed conditions, refused operation, authorization-token status, downstream commit status, protected-effect status, and no-bind receipt.

Layer 4 - Strict-Six Offline Audit Credential

OntoGuard packages proof into buyer-portable JSON, full-audit JSON/summary, governance PDF, sellable-lite PDF, decision receipt, and artifact manifest so buyers can inspect without rerunning OntoGuard.

Layer 5 - Batch DAP / Controlled Corridor Mode

The latest controlled batch proof summarizes multiple cases with aggregate decision logic, case hashes, full_audit pointers, execution-boundary summary, no-bind summary, and strict-six packaging.

Layer 6 - Semantic Projection and Learning Signals

Semantic projection, coherent governance loading, evidence geodesic, observer views, and review-gated improvement signals explain the decision without mutating live model behavior.

Boundary statement

Controlled event-level/no-bind and controlled Batch DAP proof are shown. Production L5 non-bypassability is not asserted without production integration evidence.