

OntoGuard Financial Services AI Decision Authorization

Sanitized proof case: an AI-generated investment-management recommendation was governed before reviewer-facing or customer-facing consequence.

OntoGuard governs the complete transaction through which this AI output attempted to become a business consequence.

1. Attempt COMPLETE AI-generated Investment Management portfolio construction recommendation governance withheld...	2. Interpretation COMPLETE Investment Management portfolio construction recommendation governance	3. Admissibility REVIEW REQUIRED Explains whether the proposed movement was admissible for autonomous release or required rev...	4. Authorization ESCALATE Decision API result: ESCALATE
5. Release Control WITHHELD Release authorized=False; routed_to=HUMAN_REVIEW	6. Consequence Proof NO EFFECT FORMED Autonomous release was withheld before downstream protected effect formation.	7. Human Closure PENDING Shows whether the human-review route is open, pending, or closed by an actual reviewer outcome.	8. Learning Signal CANDIDATE Buyer-visible L3 candidate only; no training, weight update, ontology mutation, policy mutat...

What was attempted	AI-generated Investment Management portfolio construction recommendation governance withheld pending governed review
What OntoGuard decided	ESCALATE - Release withheld by Decision API pending review.
What happened	WITHHELD_PENDING_REVIEW; NO_PROTECTED_EFFECT_FORMED; downstream commit=NOT_COMMITTED
What happens next	Human owner must review before release.

Proof package buyers can inspect

- Canonical governed-event lifecycle and proposed state transition.
- Execution-boundary trace and no-bind receipt.
- Decision admissibility, evidence/authority gaps, and human-review route.
- Strict-six package: buyer JSON, full-audit JSON, governance PDF, sellable-lite PDF, decision receipt, and artifact manifest.
- Benchmark disclosure: internal synthetic 2026 OntoGuard reference, 17th percentile / 50 required, n=150; not customer data or an external industry benchmark.

What this proof case is useful for

Risk / compliance	Can the proposed movement be released, withheld, blocked, or routed with a defensible record?
Model risk / governance	What evidence, uncertainty, benchmark, and authority gaps affected autonomous release?
Procurement / security	What machine-readable and buyer-readable proof exists without exposing protected internals?

Public-safe limitation

Sanitized proof case. No customer data, no protected implementation details, and no production L5 non-bypassability claim. Production route completeness requires customer route integration, endpoint enforcement evidence, bypass rejection, replay records, reviewer authorization, and outcome closure.

Contact: mark.starobinsky@ontoguard.ai | Protected by U.S. Patent Application 19/444,521 - Track I prioritized examination granted May 2026.