



# OntoGuard — Decision Authorization Infrastructure for Autonomous AI

---

## What it is (in one sentence)

OntoGuard is a **decision authorization layer** that sits **above your AI model and data** and determines whether an AI output is allowed to become a real-world action—or must be **blocked** or **escalated to a human**—with proof attached.

## The problem it solves (why executives care)

Companies want AI to drive high-stakes workflows—approvals, filings, healthcare operations, security actions, customer exceptions—but raw model output isn’t “business-admissible” because it can be wrong, unauditible, and risky.

Once AI output can trigger action, you need an **authorization + proof layer** so someone can say: “**This was allowed, here’s why, and here’s the evidence.**” OntoGuard is built to be that layer.

---

## What OntoGuard does:

Think of the model as a brilliant fast talker.

OntoGuard is the **adult in the room**:

- Treats the model’s output as a **claim, not the truth**
- Grounds it to structured enterprise reality (**objects + relationships**) and builds a **provenance-rich symbolic trace (L1)**
- Cross-checks with **semantic consensus (L2)** and then applies **alignment feedback + arbitration (L3)** to resolve conflicts and drive a single decision
- When confidence is low, it **auto-abstains / routes to review**; when confidence is high, it produces **governance-ready artifacts** that satisfy legal, risk, audit, and ops



### The 3 deliverables executives actually get (every step)

1. **Decision API (the product): ALLOW / BLOCK / ESCALATE**, plus confidence/uncertainty gating and safe abstention (no silent failures).
  2. **Evidence Pack (the moat):** sources + hashes + policy snapshot + replayable audit context (“receipts” for the decision).
  3. **Governance Report (the enterprise artifact): schema-stable JSON + PDF** with trust/risk/uncertainty, clause-level citations, fallback badges, and telemetry.
- 

### How it plugs in (enterprise reality)

OntoGuard is designed to integrate into existing stacks and run at the **decision boundary (output → action)**:

- **Inline gateway (synchronous checkpoint):** best for high-stakes flows where you must block/escalate before action
- **Async auditor (post-run governance):** best for fast pilots; still generates evidence packs + reports

And it's **LLM-agnostic and data-agnostic**: it works with your chosen model and sources—**you can swap models without rebuilding governance** (not “multiple LLMs at once” as the core promise).

---

### It also makes AI systems improve over time (training signals)

OntoGuard doesn't just police outputs. It can export **training signals** from governed decisions:

- What the AI/agent tried to do (intent/tool call)
- What evidence was missing/conflicting
- What policy/rule triggered the outcome
- What the safe/correct action should have been

This creates “clean examples” to improve RAG/agents faster—without relying on noisy internet data.



---

## The “Peak Data” angle (why it can be massively profitable)

Models go stale because new facts, policies, and edge cases (“Peak Data”) keep arriving—traditional fixes require expensive retraining and re-validation cycles.

OntoGuard is positioned as a **runtime governed mitigation layer**: update knowledge + constraints without touching base weights, enforce policy at runtime, and ship proof artifacts for audit/compliance—turning drift into a manageable (and monetizable) maintenance surface.

---

## What you license (clean expansion path)

The deck’s licensing ladder:

1. **Integration Packs**
2. **Agent Control Plane**
3. **Proof Packs**
4. **Training Signal Export**
5. **Platform Expansion**

**The simplest way to say it**

**LLMs generate intelligence. OntoGuard turns that intelligence into authorized, auditable actions—with receipts.**